



Non-math capability regression: base vs. SFT vs. RL (limit=100 per task)


LAMBADA ($\uparrow$ better)




HellaSwag ($\uparrow$ better)




ARC-Easy ($\uparrow$ better)




ARC-Challenge ($\uparrow$ better)




PIQA ($\uparrow$ better)




WinoGrande Acc ($\uparrow$ better)




OpenBookQA ($\uparrow$ better)




WikiText BPB ($\downarrow$ better)



LAMBADA Perplexity ($\downarrow$ better)



SQuAD-Completion BPB ($\downarrow$ better)



IFEval BPB ($\downarrow$ better)

