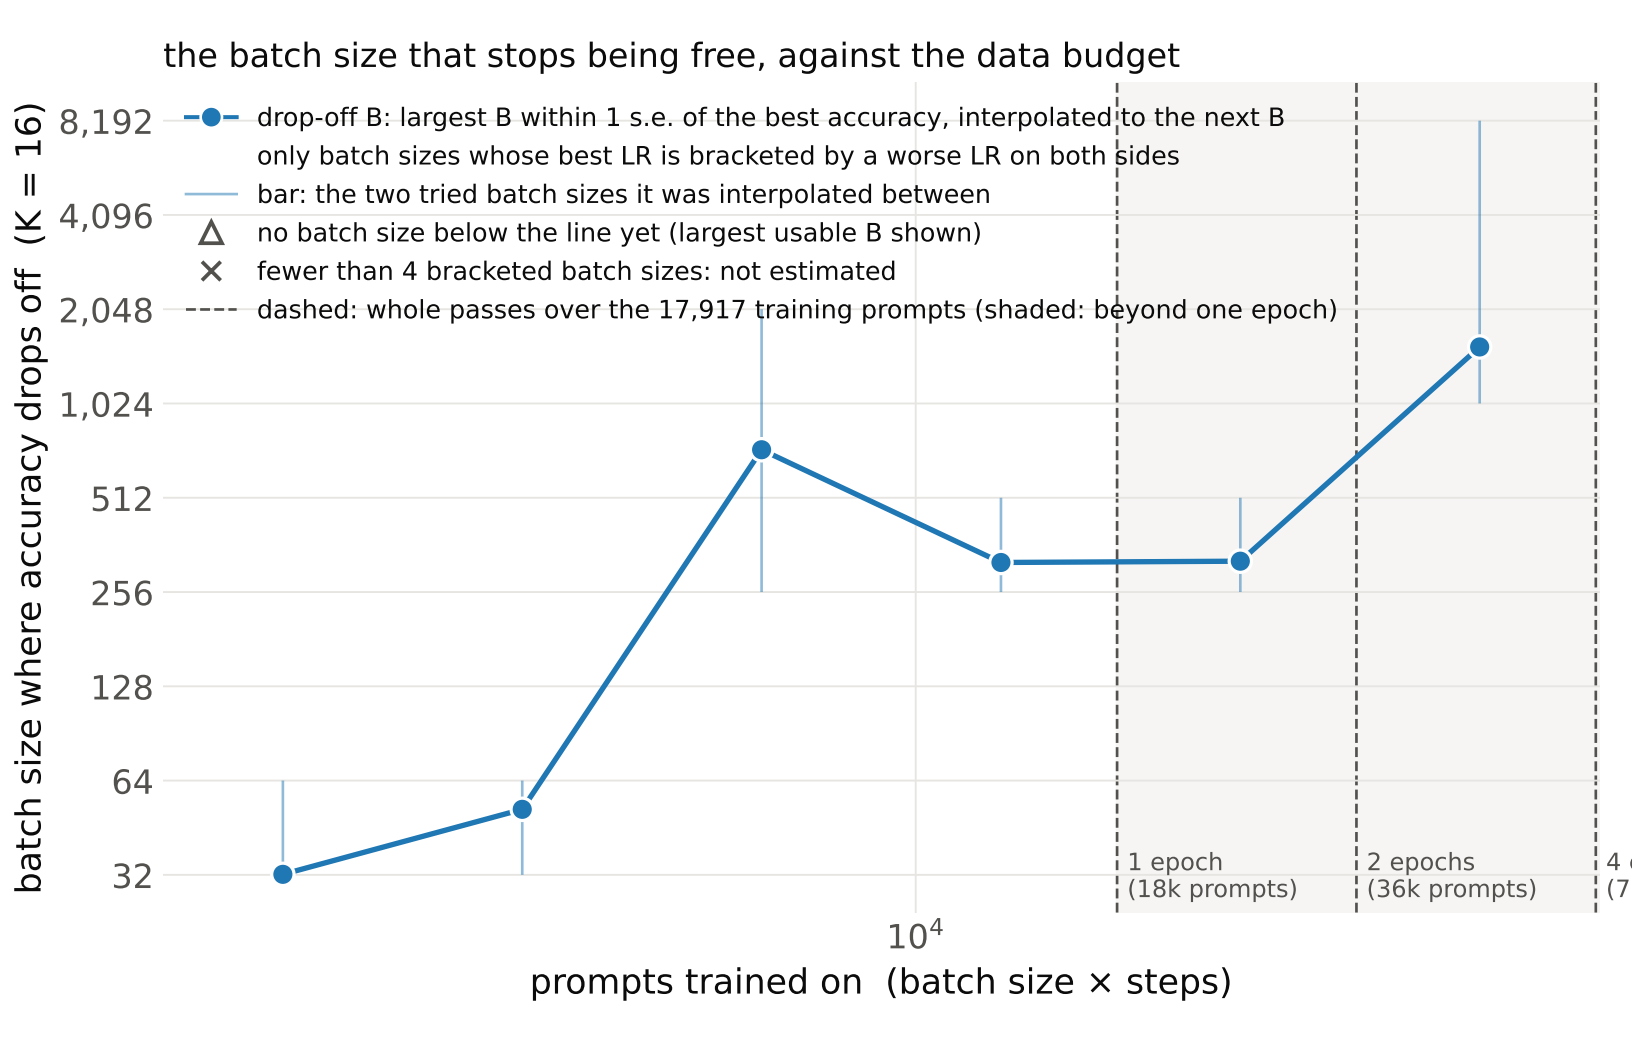
the batch size that stops being free, against the data budget
drop-off B: largest B within 1 s.e. of the best accuracy, interpolated to the next B
only batch sizes whose best LR is bracketed by a worse LR on both sides
bar: the two tried batch sizes it was interpolated between
no batch size below the line yet (largest usable B shown)
fewer than 4 bracketed batch sizes: not estimated
dashed: whole passes over the 17,917 training prompts (shaded: beyond one epoch)
batch size where accuracy drops off (K = 16)
8,192
4,096
2,048
1,024
512
256
128
64
32
1 epoch
(18k prompts)
2 epochs
(36k prompts)
10^4
prompts trained on (batch size × steps)