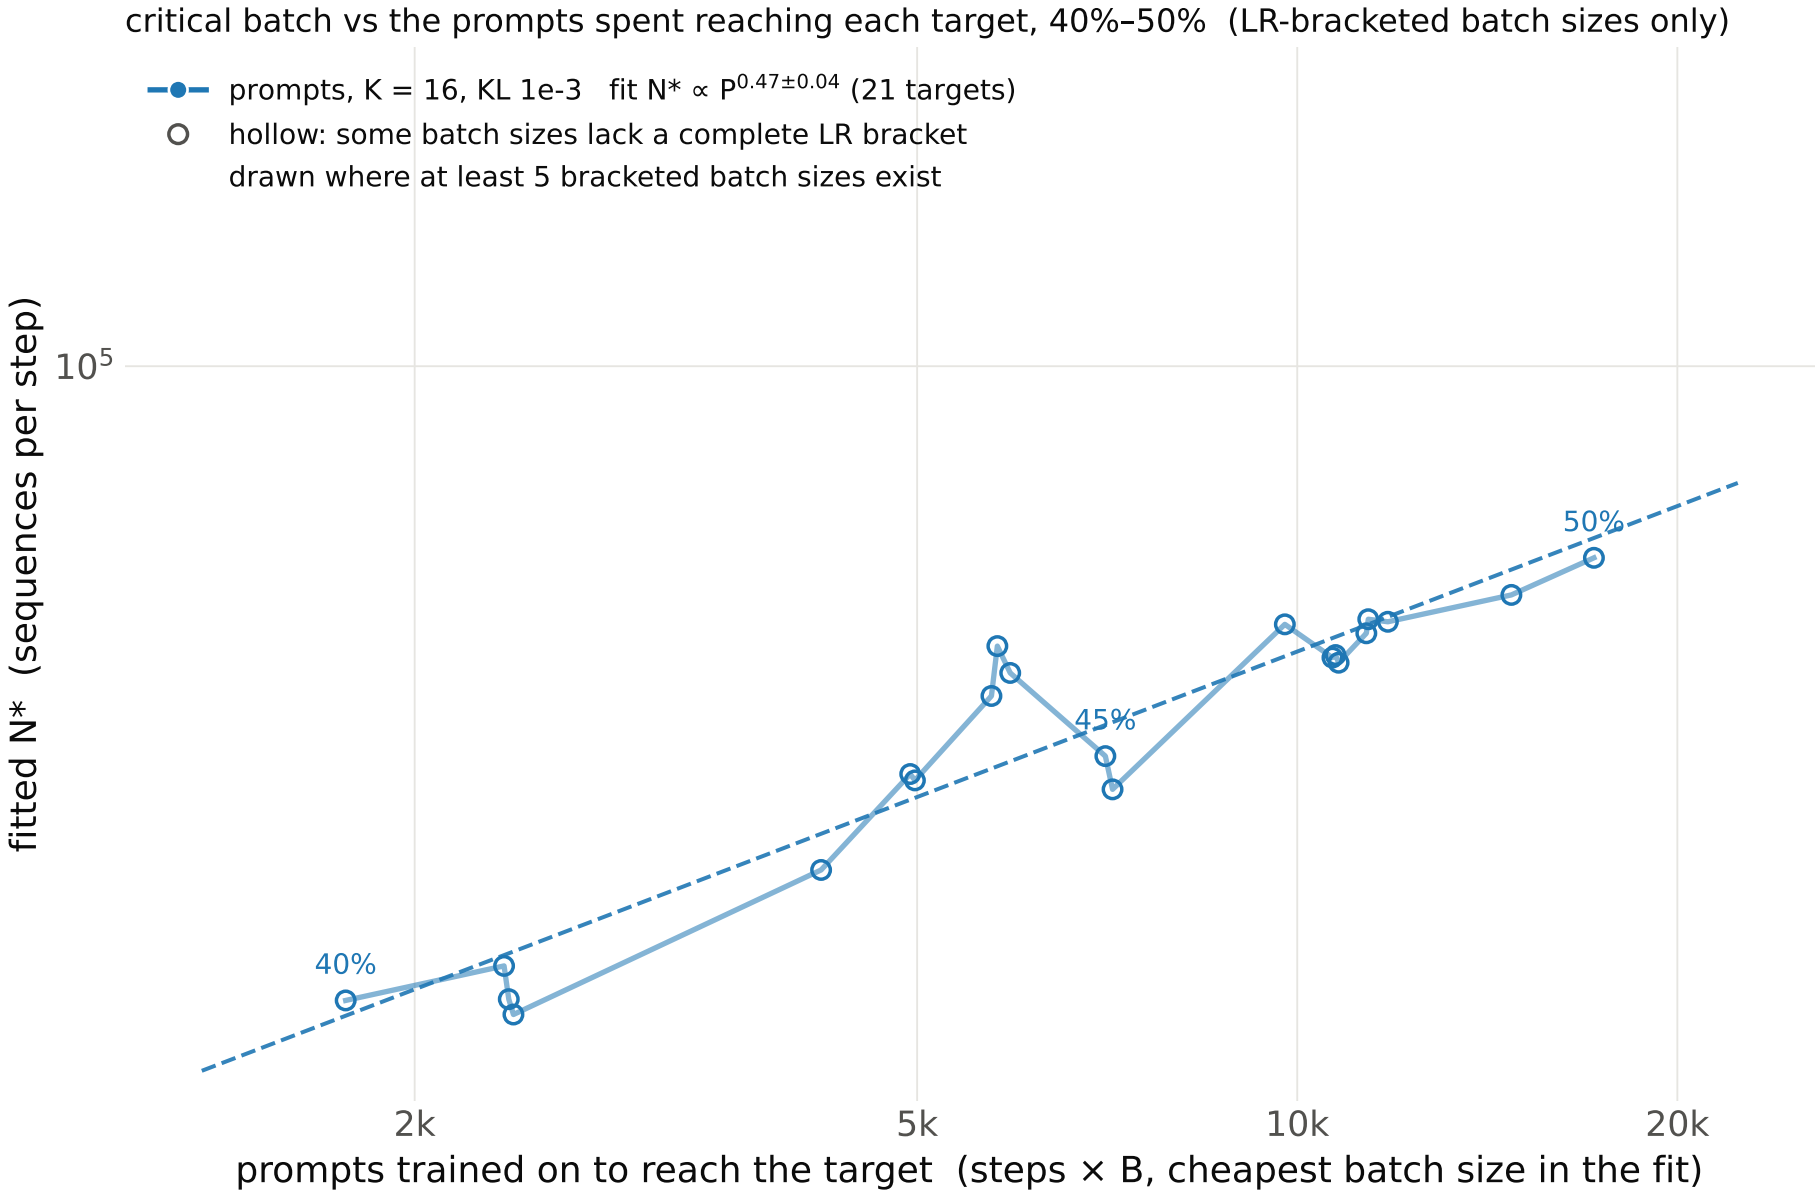
critical batch vs the prompts spent reaching each target, 40%–50% (LR-bracketed batch sizes only)
prompts, K = 16, KL 1e-3 fit N* ∝ P^0.47±0.04 (21 targets)
hollow: some batch sizes lack a complete LR bracket
drawn where at least 5 bracketed batch sizes exist
fitted N* (sequences per step)
10^5
40%
45%
50%
2k
5k
10k
20k
prompts trained on to reach the target (steps × B, cheapest batch size in the fit)