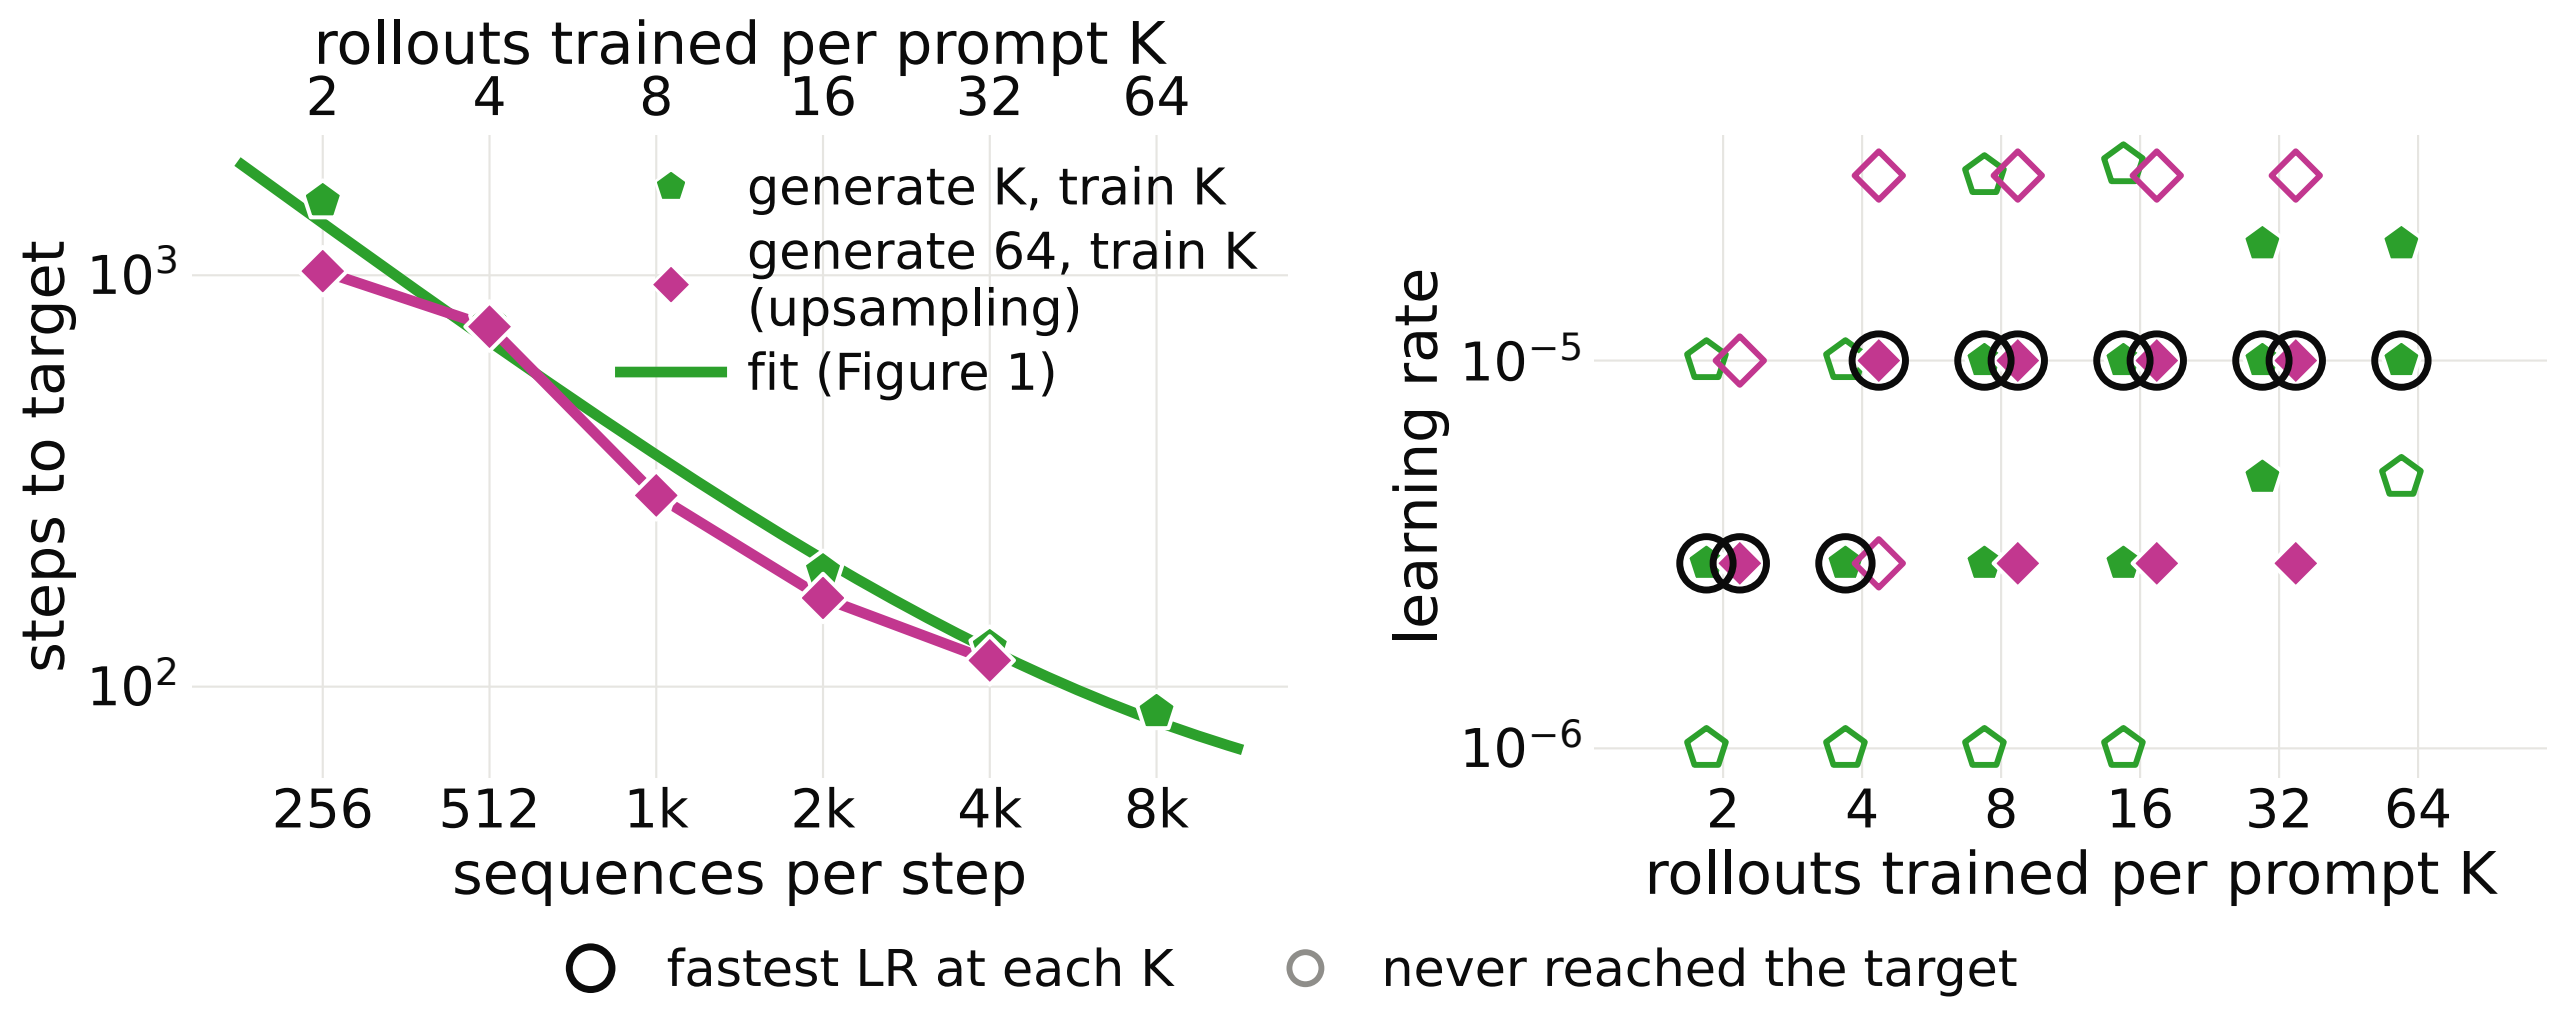
rollouts trained per prompt K
2
4
8
16
32
64
generate K, train K
generate 64, train K
(upsampling)
fit (Figure 1)
steps to target
$10^3$
$10^2$
256
512
1k
2k
4k
8k
sequences per step
learning rate
$10^{-5}$
$10^{-6}$
2
4
8
16
32
64
rollouts trained per prompt K
fastest LR at each K
never reached the target