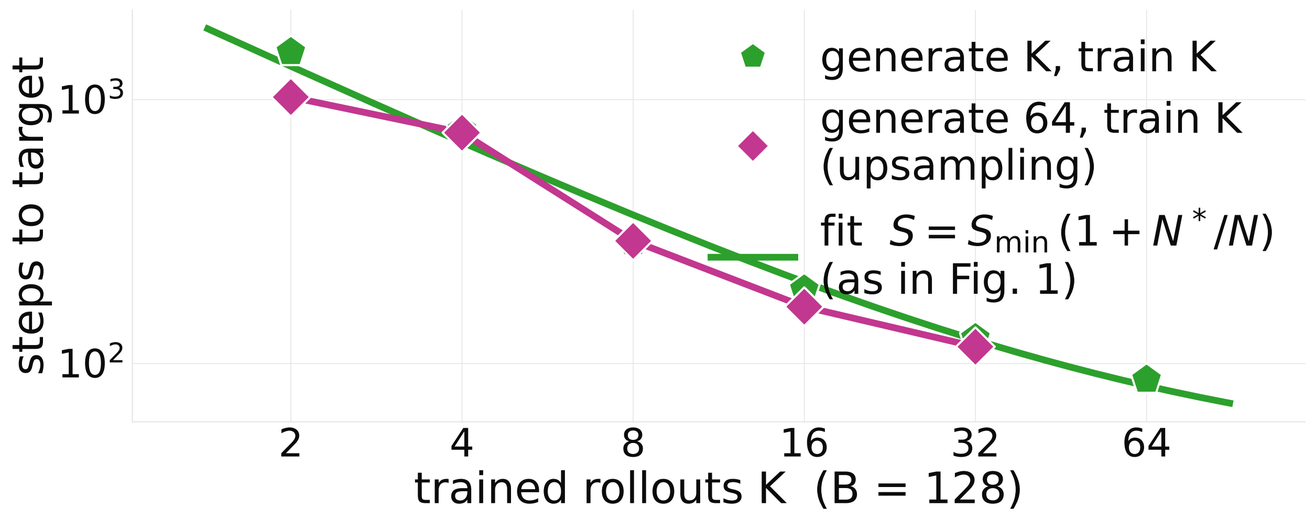
steps to target
10^3
10^2
2
4
8
16
32
64
trained rollouts K (B = 128)
generate K, train K
generate 64, train K (upsampling)
fit $S = S_{\min}(1 + N^*/N)$ (as in Fig. 1)

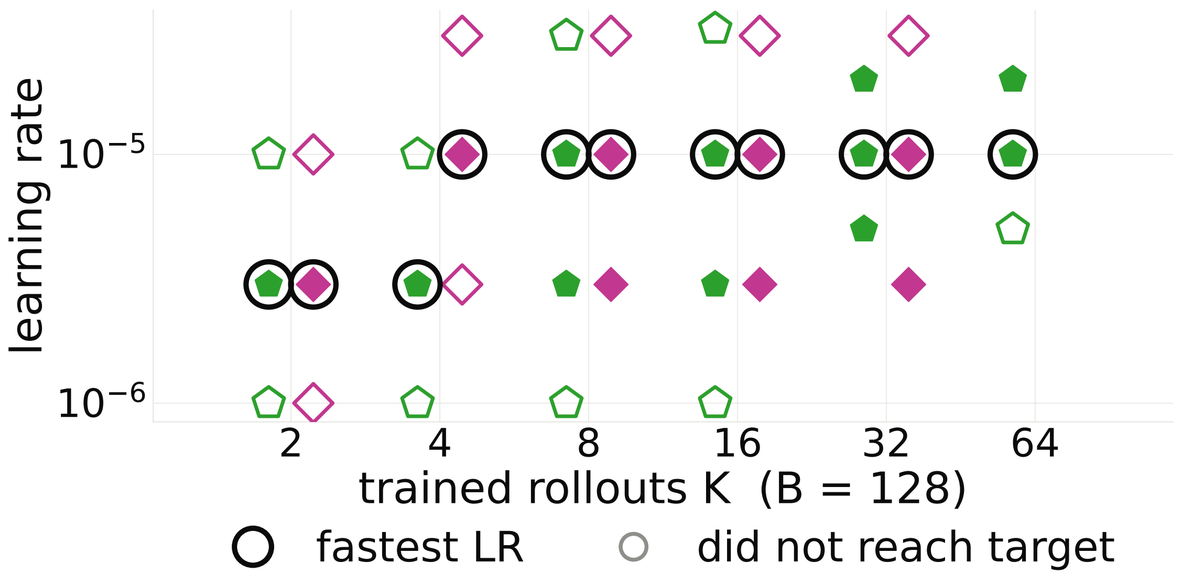
learning rate
10^{-5}
10^{-6}
2
4
8
16
32
64
trained rollouts K (B = 128)
fastest LR
did not reach target