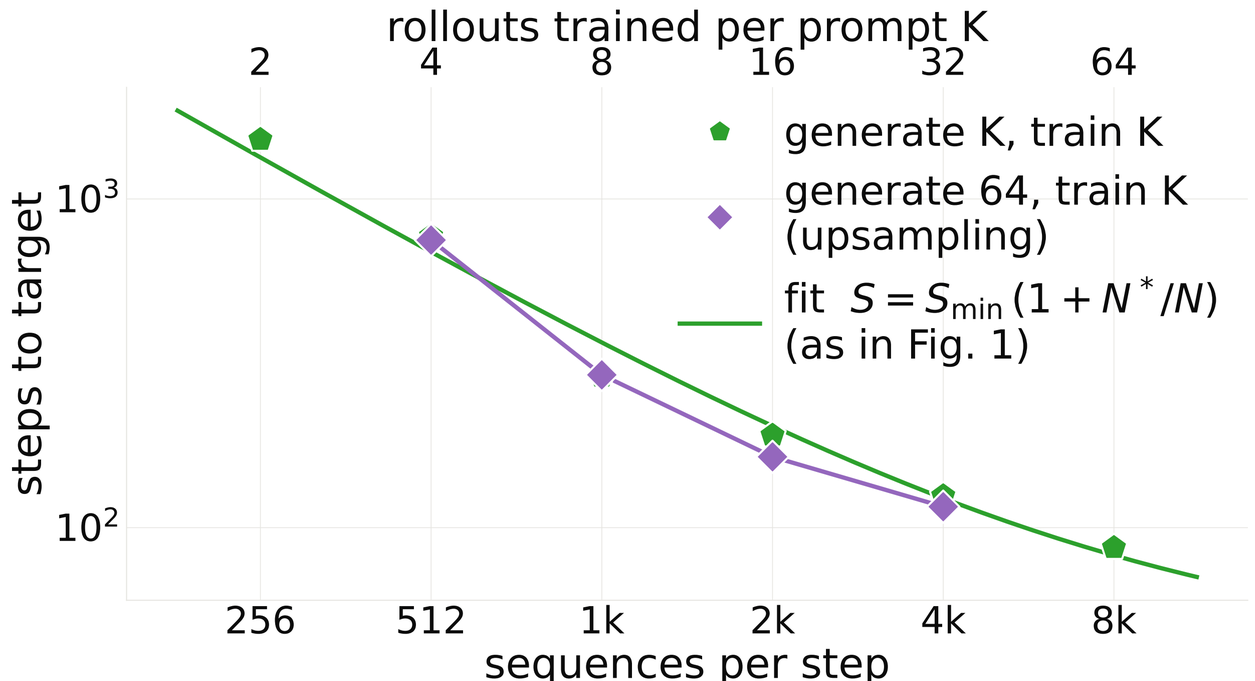
rollouts trained per prompt K
2
4
8
16
32
64
steps to target
10^3
10^2
generate K, train K
generate 64, train K
(upsampling)
fit $S = S_{\min}(1 + N^*/N)$
(as in Fig. 1)
256
512
1k
2k
4k
8k
sequences per step

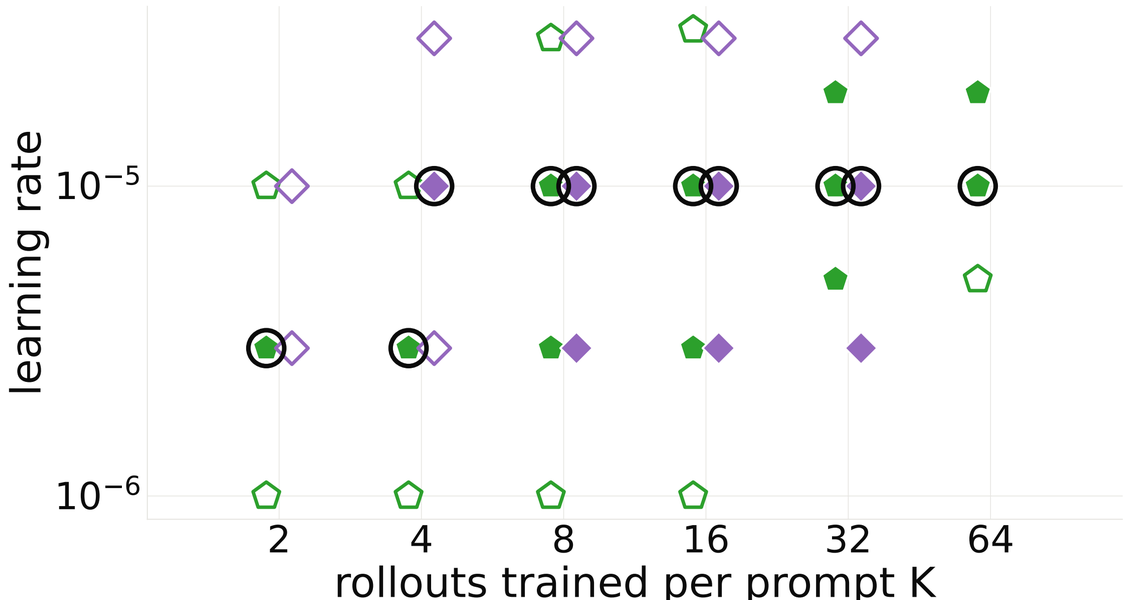
learning rate
10^{-5}
10^{-6}
2
4
8
16
32
64
rollouts trained per prompt K
fastest LR at each K
never reached the target