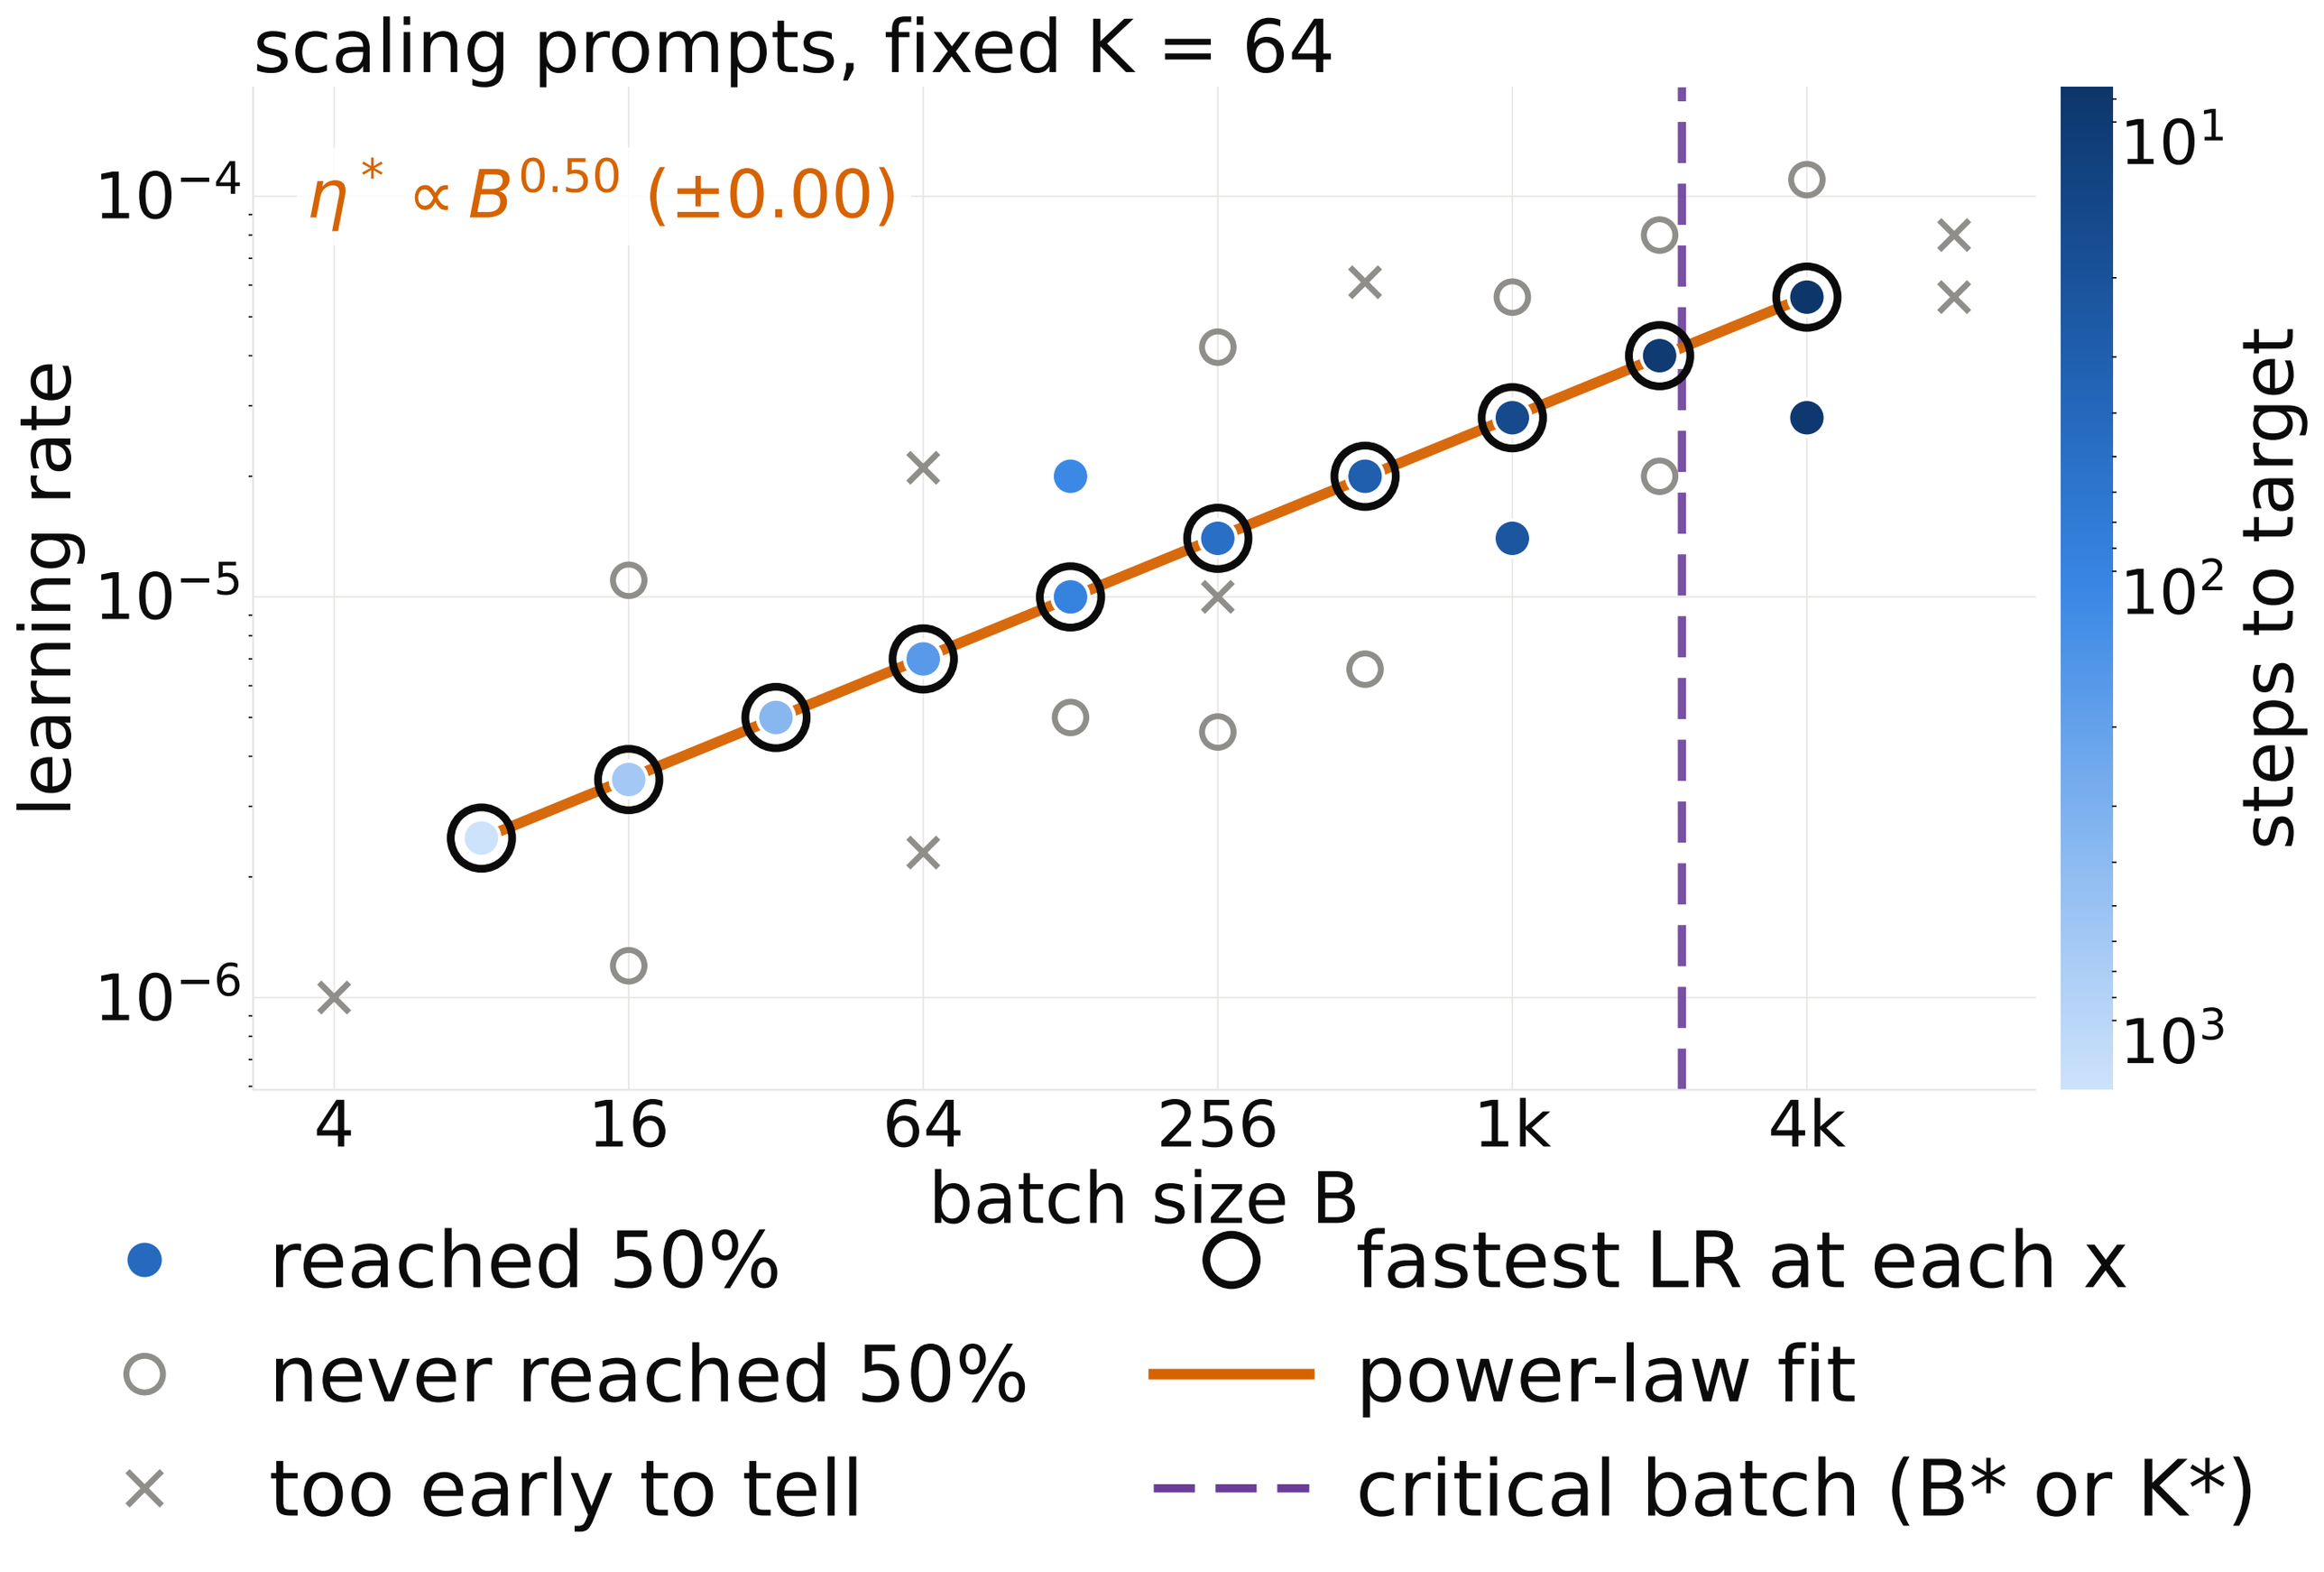
scaling prompts, fixed K = 64
$\eta^* \propto B^{0.50}$ (±0.00)
learning rate
10^{-4}
10^{-5}
10^{-6}
4
16
64
256
1k
4k
batch size B
steps to target
10^1
10^2
10^3
reached 50%
never reached 50%
too early to tell
fastest LR at each x
power-law fit
critical batch (B* or K*)